

FE590. Assignment #4.

Yifu He

2019-05-08

Instructions

When you have completed the assignment, knit the document into a PDF file, and upload *both* the .pdf and .Rmd files to Canvas.

Note that you must have LaTeX installed in order to knit the equations below. If you do not have it installed, simply delete the questions below.

```
library(knitr)
library(stringi)
library(devtools)

## Warning: package 'devtools' was built under R version 3.5.2
## Warning: package 'usethis' was built under R version 3.5.2

CWID = 10442277 #Place here your Campus wide ID number, this will personalize
#your results, but still maintain the reproduceable nature of using seeds.
#If you ever need to reset the seed in this assignment, use this as your seed
#Papers that use -1 as this CWID variable will earn 0's so make sure you change
#this value before you submit your work.
personal = CWID %% 10000
set.seed(personal)
```

Question 1:

In this assignment, you will be required to find a set of data to run regression on. This data set should be financial in nature, and of a type that will work with the models we have discussed this semester (hint: we didn't look at time series) You may not use any of the data sets in the ISLR package that we have been looking at all semester. Your data set that you choose should have both qualitative and quantitative variables. (or has variables that you can transform) Provide a description of the data below, where you obtained it, what the variable names are and what it is describing.

Answer1 :

- (i) Description of Data : The data is based on the daily historical bitcoin market prices and other factors related to it from 23th Feb 2010 to 20th Feb 2018. This data set was obtained from www.kaggle.com. The dataset includes 24 predictors and 2899 entries.

```
set.seed(personal)
setwd("/Users/yifuhe/Desktop")
data1 = read.csv("bitcoin_dataset.csv")
data = na.omit(data1)
nr = nrow(data)
nc = ncol(data)
print(paste0("The number of rows in the data set are ", nrow(data)))

## [1] "The number of rows in the data set are 2899"

print(paste0("The number of columns in the data set are ", ncol(data)))

## [1] "The number of columns in the data set are 24"

head(data)
```

##		Date	btc_market_price	btc_total_bitcoins	btc_market_cap
## 1	2010-02-23 00:00:00		0	2110700	0
## 2	2010-02-24 00:00:00		0	2120200	0
## 3	2010-02-25 00:00:00		0	2127600	0
## 4	2010-02-26 00:00:00		0	2136100	0
## 5	2010-02-27 00:00:00		0	2144750	0
## 6	2010-02-28 00:00:00		0	2152850	0
##		btc_trade_volume	btc_blocks_size	btc_avg_block_size	
## 1		0	0	0.0002163347	
## 2		0	0	0.0002817211	
## 3		0	0	0.0002269054	
## 4		0	0	0.0003186765	
## 5		0	0	0.0002234162	
## 6		0	0	0.0002914506	
##		btc_n_orphaned_blocks	btc_n_transactions_per_block		
## 1		0	1		
## 2		0	1		
## 3		0	1		
## 4		0	1		
## 5		0	1		
## 6		0	1		
##		btc_median_confirmation_time	btc_hash_rate	btc_difficulty	
## 1		0	3.153929e-05	2.527738	
## 2		0	3.571305e-05	3.781179	
## 3		0	2.781859e-05	3.781179	
## 4		0	3.195378e-05	3.781179	
## 5		0	3.251768e-05	3.781179	
## 6		0	3.045008e-05	3.781179	
##		btc_miners_revenue	btc_transaction_fees	btc_cost_per_transaction_percent	
## 1		0	0	25100.00000	

```
## 2          0          0          179.24528
## 3          0          0          1057.14286
## 4          0          0           64.58206
## 5          0          0          1922.22222
## 6          0          0          154.28571
##  btc_cost_per_transaction btc_n_unique_addresses btc_n_transactions
## 1          0          252          252
## 2          0          195          196
## 3          0          150          150
## 4          0          176          176
## 5          0          176          176
## 6          0          165          165
##  btc_n_transactions_total btc_n_transactions_excluding_popular
## 1         42613          252
## 2         42809          196
## 3         42959          150
## 4         43135          176
## 5         43311          176
## 6         43476          165
##  btc_n_transactions_excluding_chains_longer_than_100 btc_output_volume
## 1          252          12600
## 2          196          14800
## 3          150           8100
## 4          176          29349
## 5          176           9101
## 6          165          13399
##  btc_estimated_transaction_volume btc_estimated_transaction_volume_usd
## 1           50           0
## 2          5300           0
## 3           700           0
## 4         13162           0
## 5           450           0
## 6          5250           0
```

```
data$Date = as.Date(data$Date)
head(data) # after cleaning date column
```

```
##      Date btc_market_price btc_total_bitcoins btc_market_cap
## 1 2010-02-23          0          2110700          0
## 2 2010-02-24          0          2120200          0
## 3 2010-02-25          0          2127600          0
## 4 2010-02-26          0          2136100          0
## 5 2010-02-27          0          2144750          0
## 6 2010-02-28          0          2152850          0
##  btc_trade_volume btc_blocks_size btc_avg_block_size
## 1          0          0      0.0002163347
## 2          0          0      0.0002817211
## 3          0          0      0.0002269054
## 4          0          0      0.0003186765
## 5          0          0      0.0002234162
```

```

## 6          0          0      0.0002914506
##  btc_n_orphaned_blocks btc_n_transactions_per_block
## 1          0          1
## 2          0          1
## 3          0          1
## 4          0          1
## 5          0          1
## 6          0          1
##  btc_median_confirmation_time btc_hash_rate btc_difficulty
## 1          0 3.153929e-05      2.527738
## 2          0 3.571305e-05      3.781179
## 3          0 2.781859e-05      3.781179
## 4          0 3.195378e-05      3.781179
## 5          0 3.251768e-05      3.781179
## 6          0 3.045008e-05      3.781179
##  btc_miners_revenue btc_transaction_fees btc_cost_per_transaction_percent
## 1          0          0      25100.00000
## 2          0          0      179.24528
## 3          0          0     1057.14286
## 4          0          0      64.58206
## 5          0          0     1922.22222
## 6          0          0     154.28571
##  btc_cost_per_transaction btc_n_unique_addresses btc_n_transactions
## 1          0          252      252
## 2          0          195      196
## 3          0          150      150
## 4          0          176      176
## 5          0          176      176
## 6          0          165      165
##  btc_n_transactions_total btc_n_transactions_excluding_popular
## 1          42613      252
## 2          42809      196
## 3          42959      150
## 4          43135      176
## 5          43311      176
## 6          43476      165
##  btc_n_transactions_excluding_chains_longer_than_100 btc_output_volume
## 1          252      12600
## 2          196      14800
## 3          150      8100
## 4          176      29349
## 5          176      9101
## 6          165      13399
##  btc_estimated_transaction_volume btc_estimated_transaction_volume_usd
## 1          50          0
## 2         5300          0
## 3          700          0
## 4        13162          0
## 5          450          0
## 6         5250          0

```

(ii) Aim for the project : TO PREDICT BITCOIN MARKET PRICES FROM THE PREDICTORS GIVEN BELOW. ALSO TO FIND WHICH OF THE PREDICTORS ARE CORRELATED TO MARKET PRICES..

(iii) Response Variable is btc_market_price This dataset has the following features.

Date : Date of observation
btc_market_price : Average USD market price across major bitcoin exchanges.
btc_total_bitcoins : The total number of bitcoins that have already been mined.
btc_market_cap : The total USD value of bitcoin supply in circulation.
btc_trade_volume : The total USD value of trading volume on major bitcoin exchanges.
btc_blocks_size : The total size of all block headers and transactions.
btc_avg_block_size : The average block size in MB.
btc_n_orphaned_blocks : The total number of blocks mined but ultimately not attached to the main Bitcoin blockchain.
btc_n_transactions_per_block : The average number of transactions per block.
btc_median_confirmation_time : The median time for a transaction to be accepted into a mined block.
btc_hash_rate : The estimated number of tera hashes per second the Bitcoin network is performing.
btc_difficulty : A relative measure of how difficult it is to find a new block.
btc_miners_revenue : Total value of coinbase block rewards and transaction fees paid to miners.
btc_transaction_fees : The total value of all transaction fees paid to miners.
btc_cost_per_transaction_percent : miners revenue as percentage of the transaction volume.
btc_cost_per_transaction : miners revenue divided by the number of transactions.
btc_n_unique_addresses : The total number of unique addresses used on the Bitcoin blockchain.
btc_n_transactions : The number of daily confirmed Bitcoin transactions.
btc_n_transactions_total : Total number of transactions.
btc_n_transactions_excluding_popular : The total number of Bitcoin transactions, excluding the 100 most popular addresses.
btc_n_transactions_excluding_chains_longer_than_100 : The total number of Bitcoin transactions per day excluding long transaction chains.
btc_output_volume : The total value of all transaction outputs per day.
btc_estimated_transaction_volume : The total estimated value of transactions on the Bitcoin blockchain.
btc_estimated_transaction_volume_usd : The estimated transaction value in USD value.

Looking at the data....

```
set.seed(personal)
summary(data)
```

##	Date	btc_market_price	btc_total_bitcoins
##	Min. :2010-02-23	Min. : 0.000	Min. : 2110700
##	1st Qu.:2012-02-23	1st Qu.: 6.714	1st Qu.: 8410825
##	Median :2014-02-19	Median : 236.000	Median :12418575
##	Mean :2014-02-22	Mean : 901.824	Mean :11522310
##	3rd Qu.:2016-02-26	3rd Qu.: 604.460	3rd Qu.:15255538
##	Max. :2018-02-20	Max. :19498.683	Max. :16876825
##	btc_market_cap	btc_trade_volume	btc_blocks_size
##	Min. :0.000e+00	Min. :0.000e+00	Min. : 0.0
##	1st Qu.:5.488e+07	1st Qu.:2.994e+05	1st Qu.: 779.5
##	Median :3.364e+09	Median :1.024e+07	Median : 15035.0
##	Mean :1.451e+10	Mean :8.231e+07	Mean : 36202.8
##	3rd Qu.:8.229e+09	3rd Qu.:2.935e+07	3rd Qu.: 59897.5

```


## Max.      :3.265e+11    Max.      :5.352e+09    Max.      :157665.0
## btc_avg_block_size    btc_n_orphaned_blocks    btc_n_transactions_per_block
## Min.      :0.0002163    Min.      :0.0000          Min.      : 1.0
## 1st Qu.:0.0245796    1st Qu.:0.0000          1st Qu.: 54.5
## Median :0.1996229    Median :0.0000          Median : 379.0
## Mean     :0.3567457    Mean     :0.3581          Mean     : 679.3
## 3rd Qu.:0.6933442    3rd Qu.:0.0000          3rd Qu.:1245.7
## Max.      :1.1103268    Max.      :7.0000          Max.      :2722.6
## btc_median_confirmation_time    btc_hash_rate    btc_difficulty
## Min.      : 0.000          Min.      : 0          Min.      :3.000e+00
## 1st Qu.: 6.133          1st Qu.: 12          1st Qu.:1.627e+06
## Median : 7.933          Median : 25981          Median :3.130e+09
## Mean     : 7.561          Mean     : 1396897          Mean     :1.820e+11
## 3rd Qu.:10.271          3rd Qu.: 1132497          3rd Qu.:1.584e+11
## Max.      :47.733          Max.      :25579249          Max.      :2.968e+12
## btc_miners_revenue    btc_transaction_fees    btc_cost_per_transaction_percent
## Min.      : 0          Min.      : 0.000          Min.      : 0.14
## 1st Qu.: 47011          1st Qu.: 9.624          1st Qu.: 1.18
## Median : 888738          Median : 21.405          Median : 2.46
## Mean     : 2306187          Mean     : 61.201          Mean     : 58.47
## 3rd Qu.: 1862391          3rd Qu.: 51.014          3rd Qu.: 5.84
## Max.      :53191582          Max.      :1495.947          Max.      :88571.43
## btc_cost_per_transaction    btc_n_unique_addresses    btc_n_transactions
## Min.      : 0.000          Min.      : 110          Min.      : 118
## 1st Qu.: 4.172          1st Qu.: 17008          1st Qu.: 8056
## Median : 7.839          Median : 131955          Median : 62960
## Mean     : 15.192          Mean     : 196512          Mean     :103257
## 3rd Qu.: 14.976          3rd Qu.: 367857          3rd Qu.:191969
## Max.      :161.686          Max.      :1072861          Max.      :490644
## btc_n_transactions_total    btc_n_transactions_excluding_popular
## Min.      : 42613          Min.      : 118
## 1st Qu.: 2490264          1st Qu.: 6878
## Median : 33231891          Median : 54894
## Mean     : 70431859          Mean     : 95502
## 3rd Qu.:112793186          3rd Qu.:187552
## Max.      :300576632          Max.      :470650
## btc_n_transactions_excluding_chains_longer_than_100    btc_output_volume
## Min.      : 118          Min.      : 6150
## 1st Qu.: 6836          1st Qu.: 496080
## Median : 35658          Median : 1116561
## Mean     : 64000          Mean     : 1567758
## 3rd Qu.:115688          3rd Qu.: 2029856
## Max.      :318896          Max.      :45992223
## btc_estimated_transaction_volume    btc_estimated_transaction_volume_usd
## Min.      : 7          Min.      :0.000e+00
## 1st Qu.: 96478          1st Qu.:9.700e+05
## Median : 179252          Median :3.902e+07
## Mean     : 203961          Mean     :2.131e+08
## 3rd Qu.: 258903          3rd Qu.:1.386e+08
## Max.      :5825066          Max.      :5.760e+09

```

```
#Analysing the given data set:---
```

```
set.seed(personal)
```

```
plot(data$Date, data$btc_market_price, xlab = "Date", ylab = "Bitcoin_Market_Price (USD)", col = "red")
```



```
#training and testing data
```

```
t_ind = sample(1:nr, 0.75 * nr, replace = F)
```

```
train = data[t_ind,]
```

```
test = data[-t_ind,]
```

```
cor(train[, -c(1,24))][1,]
```


```
##          btc_market_price
##          1.00000000
##          btc_total_bitcoins
##          0.40648769
##          btc_market_cap
##          0.99980760
##          btc_trade_volume
##          0.86369765
##          btc_blocks_size
##          0.68955479
##          btc_avg_block_size
##          0.56025478
```

```
##          btc_n_orphaned_blocks
##          -0.08222715
##          btc_n_transactions_per_block
##          0.54527049
##          btc_median_confirmation_time
##          0.28260736
##          btc_hash_rate
##          0.90001108
##          btc_difficulty
##          0.89533355
##          btc_miners_revenue
##          0.98480128
##          btc_transaction_fees
##          0.78157994
##          btc_cost_per_transaction_percent
##          -0.01246046
##          btc_cost_per_transaction
##          0.83623012
##          btc_n_unique_addresses
##          0.66240921
##          btc_n_transactions
##          0.56562461
##          btc_n_transactions_total
##          0.69413628
##          btc_n_transactions_excluding_popular
##          0.55462051
##          btc_n_transactions_excluding_chains_longer_than_100
##          0.56676867
##          btc_output_volume
##          0.10330967
##          btc_estimated_transaction_volume
##          0.04587268
```

```
library(corrplot)
```

```
## corrplot 0.84 loaded
```

```
corrplot(cor(data[,2:12]), method = "circle")
```

(iv) After analysing the data before selecting the best predictors, some of the conclusions made are :—

- Bitcoin market prices increases exponentially wrt date but falls in between 2016 and 2018 drastically.
- From the above correlation diagram and the values calculated, predictors which seem to be correlated to the market price are : btc_market_cap, btc_trade_volume, btc_hash_rate, btc_miners_revenue, btc_difficulty, btc_cost_per_transaction

Question 2:

Pick a quantitative variable and fit at least four different models in order to predict that variable using the other predictors. Determine which of the models is the best fit. You will need to provide strong reasons as to why the particular model you chose is the best one. You will need to confirm the model you have selected provides the best fit and that you have obtained the best version of that particular model (i.e. subset selection or validation for example). You need to convince the grader that you have chosen the best model.

Answer 2 :

(Model 1) SIMPLE LINEAR REGRESSION : Since this is a simple regression model .. let us select the best predictor instead of applying to all the predictors. Also after finding correlation, let us see what are the results after the best subset selection, forward subset selection and backward subset selection.

```
set.seed(personal)
library(leaps)
p = regsubsets(btc_market_price ~ btc_market_cap + btc_trade_volume + btc_hash_rate + btc_miners_revenue + btc_difficulty + btc_cost_per_transaction, data = train)
q = regsubsets(btc_market_price ~ btc_market_cap + btc_trade_volume + btc_hash_rate + btc_miners_revenue + btc_difficulty + btc_cost_per_transaction, data = train, method = "forward")
r = regsubsets(btc_market_price ~ btc_market_cap + btc_trade_volume + btc_hash_rate + btc_miners_revenue + btc_difficulty + btc_cost_per_transaction, data = train, method = "backward")
summary(p)[7]

## $outmat
##          btc_market_cap btc_trade_volume btc_hash_rate btc_miners_revenue
## 1 ( 1 ) "*"           " "               " "             " "
## 2 ( 1 ) "*"           " "               " "             " "
## 3 ( 1 ) "*"           " "               " "             "*"
## 4 ( 1 ) "*"           "*"              " "             "*"
## 5 ( 1 ) "*"           " "               "*"             "*"
## 6 ( 1 ) "*"           "*"              "*"             "*"
##          btc_difficulty btc_cost_per_transaction
## 1 ( 1 ) " "           " "
## 2 ( 1 ) " "           "*"
## 3 ( 1 ) " "           "*"
## 4 ( 1 ) " "           "*"
## 5 ( 1 ) "*"           "*"
## 6 ( 1 ) "*"           "*"

summary(q)[7]

## $outmat
##          btc_market_cap btc_trade_volume btc_hash_rate btc_miners_revenue
## 1 ( 1 ) "*"           " "               " "             " "
## 2 ( 1 ) "*"           " "               " "             " "
## 3 ( 1 ) "*"           " "               " "             "*"
## 4 ( 1 ) "*"           "*"              " "             "*"
## 5 ( 1 ) "*"           "*"              "*"             "*"
## 6 ( 1 ) "*"           "*"              "*"             "*"
##          btc_difficulty btc_cost_per_transaction
## 1 ( 1 ) " "           " "
## 2 ( 1 ) " "           "*"
## 3 ( 1 ) " "           "*"
## 4 ( 1 ) " "           "*"
## 5 ( 1 ) " "           "*"
## 6 ( 1 ) " "           "*"
## 7 ( 1 ) " "           " "
## 8 ( 1 ) " "           " "
## 9 ( 1 ) " "           " "
## 10 ( 1 ) " "          " "
## 11 ( 1 ) " "          " "
## 12 ( 1 ) " "          " "
## 13 ( 1 ) " "          " "
## 14 ( 1 ) " "          " "
## 15 ( 1 ) " "          " "
## 16 ( 1 ) " "          " "
## 17 ( 1 ) " "          " "
## 18 ( 1 ) " "          " "
## 19 ( 1 ) " "          " "
## 20 ( 1 ) " "          " "
## 21 ( 1 ) " "          " "
## 22 ( 1 ) " "          " "
## 23 ( 1 ) " "          " "
## 24 ( 1 ) " "          " "
## 25 ( 1 ) " "          " "
## 26 ( 1 ) " "          " "
## 27 ( 1 ) " "          " "
## 28 ( 1 ) " "          " "
## 29 ( 1 ) " "          " "
## 30 ( 1 ) " "          " "
## 31 ( 1 ) " "          " "
## 32 ( 1 ) " "          " "
## 33 ( 1 ) " "          " "
## 34 ( 1 ) " "          " "
## 35 ( 1 ) " "          " "
## 36 ( 1 ) " "          " "
## 37 ( 1 ) " "          " "
## 38 ( 1 ) " "          " "
## 39 ( 1 ) " "          " "
## 40 ( 1 ) " "          " "
## 41 ( 1 ) " "          " "
## 42 ( 1 ) " "          " "
## 43 ( 1 ) " "          " "
## 44 ( 1 ) " "          " "
## 45 ( 1 ) " "          " "
## 46 ( 1 ) " "          " "
## 47 ( 1 ) " "          " "
## 48 ( 1 ) " "          " "
## 49 ( 1 ) " "          " "
## 50 ( 1 ) " "          " "
## 51 ( 1 ) " "          " "
## 52 ( 1 ) " "          " "
## 53 ( 1 ) " "          " "
## 54 ( 1 ) " "          " "
## 55 ( 1 ) " "          " "
## 56 ( 1 ) " "          " "
## 57 ( 1 ) " "          " "
## 58 ( 1 ) " "          " "
## 59 ( 1 ) " "          " "
## 60 ( 1 ) " "          " "
## 61 ( 1 ) " "          " "
## 62 ( 1 ) " "          " "
## 63 ( 1 ) " "          " "
## 64 ( 1 ) " "          " "
## 65 ( 1 ) " "          " "
## 66 ( 1 ) " "          " "
## 67 ( 1 ) " "          " "
## 68 ( 1 ) " "          " "
## 69 ( 1 ) " "          " "
## 70 ( 1 ) " "          " "
## 71 ( 1 ) " "          " "
## 72 ( 1 ) " "          " "
## 73 ( 1 ) " "          " "
## 74 ( 1 ) " "          " "
## 75 ( 1 ) " "          " "
## 76 ( 1 ) " "          " "
## 77 ( 1 ) " "          " "
## 78 ( 1 ) " "          " "
## 79 ( 1 ) " "          " "
## 80 ( 1 ) " "          " "
## 81 ( 1 ) " "          " "
## 82 ( 1 ) " "          " "
## 83 ( 1 ) " "          " "
## 84 ( 1 ) " "          " "
## 85 ( 1 ) " "          " "
## 86 ( 1 ) " "          " "
## 87 ( 1 ) " "          " "
## 88 ( 1 ) " "          " "
## 89 ( 1 ) " "          " "
## 90 ( 1 ) " "          " "
## 91 ( 1 ) " "          " "
## 92 ( 1 ) " "          " "
## 93 ( 1 ) " "          " "
## 94 ( 1 ) " "          " "
## 95 ( 1 ) " "          " "
## 96 ( 1 ) " "          " "
## 97 ( 1 ) " "          " "
## 98 ( 1 ) " "          " "
## 99 ( 1 ) " "          " "
## 100 ( 1 ) " "         " "
```

```
## 4 ( 1 ) " " "*"
## 5 ( 1 ) " " "*"
## 6 ( 1 ) "*" "*"

summary(r)[7]

## $outmat
##          btc_market_cap btc_trade_volume btc_hash_rate btc_miners_revenue
## 1 ( 1 ) "*"          " "          " "          " "
## 2 ( 1 ) "*"          " "          " "          " "
## 3 ( 1 ) "*"          " "          " "          "*"
## 4 ( 1 ) "*"          " "          "*"          "*"
## 5 ( 1 ) "*"          " "          "*"          "*"
## 6 ( 1 ) "*"          "*"          "*"          "*"
##          btc_difficulty btc_cost_per_transaction
## 1 ( 1 ) " "          " "
## 2 ( 1 ) " "          "*"
## 3 ( 1 ) " "          "*"
## 4 ( 1 ) " "          "*"
## 5 ( 1 ) "*"          "*"
## 6 ( 1 ) "*"          "*"

```

After seeing the best, forward and the backward subset selection, the `btc_market_cap` i.e the total USD of bitcoin in circulation is the best predictor. So applying simple linear regression using this predictor.

```
set.seed(personal)
model1 = lm(btc_market_price ~ btc_market_cap, data = train)
model11 = lm(btc_market_price ~ btc_cost_per_transaction, data = train)
summary(model1)

##
## Call:
## lm(formula = btc_market_price ~ btc_market_cap, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -96.264 -34.972  -4.253   7.860 285.249
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.791e+01  1.093e+00   34.7   <2e-16 ***
## btc_market_cap 5.956e-08  2.507e-11  2375.5   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 47.97 on 2172 degrees of freedom
## Multiple R-squared:  0.9996, Adjusted R-squared:  0.9996
## F-statistic: 5.643e+06 on 1 and 2172 DF, p-value: < 2.2e-16

summary(model11)

```

```
##
## Call:
## lm(formula = btc_market_price ~ btc_cost_per_transaction, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -7196.7  -254.7   153.0   527.6  7766.0
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -527.598     35.182  -15.00  <2e-16 ***
## btc_cost_per_transaction    93.914      1.321   71.07  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1341 on 2172 degrees of freedom
## Multiple R-squared:  0.6993, Adjusted R-squared:  0.6991
## F-statistic: 5051 on 1 and 2172 DF, p-value: < 2.2e-16
```

We can see how strongly `btc_market_cap` affects the `btc_market` price as $\text{adj } R^2$ is nearly equal to 1. Also the second best predictor `btc_cost_per_transaction` has a adjusted R-squared = 0.7012. Both are statistically significant as seen by the p-values. Now we test this and predict values using testing dataset.

```
set.seed(personal)
pred1 = predict(model1, newdata = test)
pred11 = predict(model11, newdata = test)
mss1 = mean((test$btc_market_price-pred1)^2)
mss11 = mean((test$btc_market_price-pred11)^2)
mss1

## [1] 2366.149

mss11


## [1] 1730603
```

MSS is too high so this is not a good model. Simple Linear regression is not good to fit. So we proceed to multiple regression.

(Model ii) Multiple Linear Regression :

```
c1 <- summary(p)$cp
plot(c1,type='b',xlab="No. of Predictors",ylab=expression("Mallows C"[P]), col="red")

points(which.min(c1), c1[which.min(c1)], pch=20, col="red")
```



All the 6 predictors we selected have the minimum mallows Cp. So applying multiple regression model. But since btc_market_cap has adj R² = 0.996, we perform multiple linear regression with it and one multiple linear regression without it to exclude other variables as it is overshadowing others.

```
model12 = lm(btc_market_price ~ btc_market_cap + btc_trade_volume + btc_hash_rate + btc_miners_revenue + btc_difficulty + btc_cost_per_transaction, data = train)
summary(model12)
```

```
##
## Call:
## lm(formula = btc_market_price ~ btc_market_cap + btc_trade_volume +
##     btc_hash_rate + btc_miners_revenue + btc_difficulty + btc_cost_per_transaction,
##     data = train)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-173.439	-13.540	-2.434	15.228	178.729

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	9.427e+00	8.232e-01	11.452	< 2e-16 ***

```
## btc_market_cap          5.622e-08  1.727e-10 325.458 < 2e-16 ***
## btc_trade_volume        -2.816e-08  4.062e-09  -6.931 5.49e-12 ***
## btc_hash_rate           -2.151e-05  1.576e-06 -13.649 < 2e-16 ***
## btc_miners_revenue       1.549e-05  1.093e-06  14.172 < 2e-16 ***
## btc_difficulty           1.880e-10  1.431e-11  13.137 < 2e-16 ***
## btc_cost_per_transaction 2.572e+00  5.970e-02  43.088 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 28.34 on 2167 degrees of freedom
## Multiple R-squared:  0.9999, Adjusted R-squared:  0.9999
## F-statistic: 2.695e+06 on 6 and 2167 DF,  p-value: < 2.2e-16

model13 = lm(btc_market_price ~ btc_trade_volume + btc_hash_rate + btc_miners
_revenue + btc_difficulty + btc_cost_per_transaction, data = train)
summary(model13)

##
## Call:
## lm(formula = btc_market_price ~ btc_trade_volume + btc_hash_rate +
##     btc_miners_revenue + btc_difficulty + btc_cost_per_transaction,
##     data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3432.5   -54.4    47.2    73.8   1636.1
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -6.371e+01  5.592e+00 -11.394 < 2e-16 ***
## btc_trade_volume -2.327e-07  2.834e-08  -8.213 3.68e-16 ***
## btc_hash_rate   -2.781e-04  9.633e-06 -28.873 < 2e-16 ***
## btc_miners_revenue  3.558e-04  2.256e-06 157.734 < 2e-16 ***
## btc_difficulty   3.464e-09  7.186e-11  48.199 < 2e-16 ***
## btc_cost_per_transaction -5.133e+00  3.869e-01 -13.264 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 200.1 on 2168 degrees of freedom
## Multiple R-squared:  0.9933, Adjusted R-squared:  0.9933
## F-statistic: 6.444e+04 on 5 and 2168 DF,  p-value: < 2.2e-16
```

So after applying multiple linear regression without the `btc_market_cap`, the conclusion is that all the other variables too are significant.

```
set.seed(personal)
pred2 = predict(model12, newdata = test)
mss2 = mean((test$btc_market_price - pred2)^2)
mss2

## [1] 887.268
```

Still the error is significantly high but lower than simple linear. So we proceed to splines

(Model 3) : Polynomial regression : By the subset selections now select only and the above results drop btc_trade_volume

```
set.seed(personal)
fit1 = lm(btc_market_price ~ poly(btc_market_cap, 2), data = train)
fit2 = lm(btc_market_price ~ poly(btc_market_cap, 3), data = train)
summary(fit1)

##
## Call:
## lm(formula = btc_market_price ~ poly(btc_market_cap, 2), data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -106.601  -31.645   -6.336    4.144   281.421
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      912.059      1.014  899.600  < 2e-16 ***
## poly(btc_market_cap, 2)1 113959.971    47.272 2410.731  < 2e-16 ***
## poly(btc_market_cap, 2)2  -383.997    47.272   -8.123 7.54e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 47.27 on 2171 degrees of freedom
## Multiple R-squared:  0.9996, Adjusted R-squared:  0.9996
## F-statistic: 2.906e+06 on 2 and 2171 DF, p-value: < 2.2e-16

summary(fit2)

##
## Call:
## lm(formula = btc_market_price ~ poly(btc_market_cap, 3), data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -135.140  -25.472  -17.719    4.841   268.138
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    9.121e+02  9.546e-01  955.389  <2e-16 ***
## poly(btc_market_cap, 3)1  1.140e+05  4.451e+01 2560.235  <2e-16 ***
## poly(btc_market_cap, 3)2 -3.840e+02  4.451e+01   -8.627  <2e-16 ***
## poly(btc_market_cap, 3)3  7.430e+02  4.451e+01   16.692  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 44.51 on 2170 degrees of freedom
```

```

## Multiple R-squared:  0.9997, Adjusted R-squared:  0.9997
## F-statistic: 2.185e+06 on 3 and 2170 DF,  p-value: < 2.2e-16

fita = lm(btc_market_price ~ btc_trade_volume + btc_hash_rate + btc_miners_revenue + btc_difficulty + btc_cost_per_transaction + poly(btc_market_cap, 2), data = train)
fitb = lm(btc_market_price ~ btc_trade_volume + btc_hash_rate + btc_miners_revenue + btc_difficulty + btc_cost_per_transaction + poly(btc_market_cap, 3), data = train)
fitc = lm(btc_market_price ~ btc_trade_volume + btc_hash_rate + btc_miners_revenue + btc_difficulty + btc_cost_per_transaction + poly(btc_market_cap, 4), data = train)
fitd = lm(btc_market_price ~ btc_trade_volume + btc_hash_rate + btc_miners_revenue + btc_difficulty + btc_cost_per_transaction + poly(btc_market_cap, 5), data = train)
anova(fita, fitb, fitc, fitd)

## Analysis of Variance Table
##
## Model 1: btc_market_price ~ btc_trade_volume + btc_hash_rate + btc_miners_revenue +
##      btc_difficulty + btc_cost_per_transaction + poly(btc_market_cap,
##      2)
## Model 2: btc_market_price ~ btc_trade_volume + btc_hash_rate + btc_miners_revenue +
##      btc_difficulty + btc_cost_per_transaction + poly(btc_market_cap,
##      3)
## Model 3: btc_market_price ~ btc_trade_volume + btc_hash_rate + btc_miners_revenue +
##      btc_difficulty + btc_cost_per_transaction + poly(btc_market_cap,
##      4)
## Model 4: btc_market_price ~ btc_trade_volume + btc_hash_rate + btc_miners_revenue +
##      btc_difficulty + btc_cost_per_transaction + poly(btc_market_cap,
##      5)
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1    2166 1024462
## 2    2165  690414  1    334048 1318.18 < 2.2e-16 ***
## 3    2164  650834  1     39580  156.18 < 2.2e-16 ***
## 4    2163  548141  1    102694  405.24 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

pred3 = predict(fit1, newdata = test)
mss3 = mean((test$btc_market_cap - pred3)^2)
mss3

## [1] 1.67828e+21

```

The mse for this is higher than simple and multiple linear regression. So we discard this model.

(Model iv) So this time I choose random forest over GAM and other techniques such as splines because they all have polynomials involved and by intuition I thought those would give somewhat the same result. The reason I chose the random forest dataset size is medium not large

```
set.seed(personal)
library(randomForest)

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

?randomForest
model4 <- randomForest(x = train[,c(4, 10, 11, 12, 15)], y = train$btc_market_price, ntree = 501)
summary(model4)

##               Length Class  Mode
## call              4  -none-  call
## type              1  -none- character
## predicted        2174  -none-  numeric
## mse              501  -none-  numeric
## rsq              501  -none-  numeric
## oob.times        2174  -none-  numeric
## importance         5  -none-  numeric
## importanceSD       0  -none-  NULL
## localImportance    0  -none-  NULL
## proximity          0  -none-  NULL
## ntree             1  -none-  numeric
## mtry              1  -none-  numeric
## forest            11  -none-  list
## coefs              0  -none-  NULL
## y                2174  -none-  numeric
## test              0  -none-  NULL
## inbag             0  -none-  NULL

pred4 = predict(model4, newdata = test)
mss4 = mean((test$btc_market_price - pred4)^2)
mss4

## [1] 35380.11
```

SO my intuition that the random forest will produce better results on test is false. Hence i Will choose multiple linear regression over the other models as it has the less test error. Also this is because of overfitting. All the models fit into the training set but do poorly on test data. This can be expected as this is a to predict future bitcoin market prices.

Question 3:

Do the same approach as in question 2, but this time for a qualitative variable.

Answer 3 : As there is no binary variable for the above dataset lets create one :—

```
set.seed(personal)
mprice = mean(data$btc_market_price)
mprice

## [1] 901.8236

market_price <- rep("Greaterthanmean", nr)
market_price[data$btc_market_price < mprice] <- "Not_greater_than_mean"
data$market_price <- market_price
data$market_price <- as.factor(data$market_price)
```

(Model i) : Logistic Regression Model : Now Im changing the test set and then taking the 6 predictors for which my correlation was high.

```
set.seed(personal)
t_ind2 = sample(1:nr, 0.70 * nr, replace = FALSE)
train2 = data[t_ind2,]
test2 = data[-t_ind2,]
direction = test2$market_price
glm.fit = glm(market_price ~ btc_market_cap + btc_trade_volume + btc_hash_rate + btc_miners_revenue + btc_difficulty + btc_cost_per_transaction, data = train2, family = binomial)

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

prob = predict(glm.fit, test2, type = "response")
nrtest2 = round(0.30 * nr)
p11 = rep("Not_greater_than_mean", nrtest2)
p11[prob > 0.5] = "Greaterthanmean"
table(p11, direction)

##                direction
## p11                Greaterthanmean Not_greater_than_mean
## Greaterthanmean                2                760
## Not_greater_than_mean            108                0
```

No need to calculate further as Logistic Regression performed very badly on this dataset.

(Model ii) LDA Model:—

```
library(MASS)
set.seed(personal)
fit33 = lda(market_price ~ btc_market_cap + btc_trade_volume + btc_hash_rate + btc_miners_revenue + btc_difficulty + btc_cost_per_transaction, data = train2)
p2 = predict(fit33, test2)$class
table(p2, direction)
```

```
##                                direction
## p2                Greaterthanmean Not_greater_than_mean
## Greaterthanmean                86                0
## Not_greater_than_mean            24               760

#Misclassification Error rate
misc1 = mean(p2 != direction)
misc1

## [1] 0.02758621
```

We see only 35 datapoints were misclassified. LDA gave a very low missclassification error.

(Model iii)

```
set.seed(personal)
fit44 = qda(market_price ~ btc_market_cap + btc_trade_volume + btc_hash_rate
+ btc_miners_revenue + btc_difficulty + btc_cost_per_transaction, data = trai
n2)
p3 = predict(fit44, test2)$class
table(p3, direction)

##                                direction
## p3                Greaterthanmean Not_greater_than_mean
## Greaterthanmean                107                17
## Not_greater_than_mean            3               743

#Misclassification Error rate
misc11 = mean(p3 != direction)
misc11

## [1] 0.02298851
```

QDA performs better than LDA and gives an error rate of 0.0206 which is considered to be very low.

(Model iv) KNN :


```
set.seed(personal)
library(class)
trainKNN = as.matrix(data.frame(train2$btc_market_cap, train2$btc_trade_volum
e, train2$btc_miners_revenue, train2$btc_difficulty, train2$btc_cost_per_tran
saction))
testKNN = as.matrix(data.frame(test2$btc_market_cap, test2$btc_trade_volume,
test2$btc_miners_revenue, test2$btc_difficulty, test2$btc_cost_per_transactio
n))
direction2 = train2$market_price
pKNN1 = knn(trainKNN, testKNN, direction2, k=1)
pKNN2 = knn(trainKNN, testKNN, direction2, k=3)
pKNN3 = knn(trainKNN, testKNN, direction2, k=5)
pKNN4 = knn(trainKNN, testKNN, direction2, k=10)
pKNN5 = knn(trainKNN, testKNN, direction2, k=25)
```

```

pKNN6 = knn(trainKNN, testKNN, direction2,k=67)

error1 = mean(pKNN1 != direction)
error2 = mean(pKNN2 != direction)
error3 = mean(pKNN3 != direction)
error4 = mean(pKNN4 != direction)
error5 = mean(pKNN5 != direction)
error6 = mean(pKNN6 != direction)
error = c(error1, error2, error3, error4, error5, error6)
plot(error, type='b', xlab="k", ylab="Error", col="red")
points(which.min(error), error[which.min(error)], pch=20, col="black")

```



```
error1
```

```
## [1] 0.002298851
```

Out of all the models KNN gives the lowest missclassification error on the test set . Hence KNN will be selected to give the best prediction for the data and the created binary variable #Question 4:

(Based on ISLR Chapter 9 #7) In this problem, you will use support vector approaches in order to predict whether a given car gets high or low gas mileage based on the Auto data set.

(a)

Create a binary variable that takes on a 1 for cars with gas mileage above the median, and a 0 for cars with gas mileage below the median.

```
set.seed(personal)
library(ISLR)
attach(Auto)
#To make a binary variable use ifelse
median_mileage = median(Auto$mpg)
bin.var = ifelse(Auto$mpg > median_mileage, 1, 0)
Auto$binary.mpg = as.factor(bin.var)
```

(b)

Fit a support vector classifier to the data with various values of cost, in order to predict whether a car gets high or low gas mileage. Report the cross-validation errors associated with different values of this parameter. Comment on your results.

```
set.seed(personal)
library(e1071)

## Warning: package 'e1071' was built under R version 3.5.2

x = tune(svm, binary.mpg~., data = Auto, kernel = "linear", ranges = list(cost = c(0.01, 0.1, 1, 5, 10, 100)))
summary(x)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     1
##
## - best performance: 0.01275641
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-02 0.07391026 0.04245856
## 2 1e-01 0.05108974 0.04191745
## 3 1e+00 0.01275641 0.01344780
## 4 5e+00 0.01782051 0.01229997
## 5 1e+01 0.01782051 0.01229997
## 6 1e+02 0.03051282 0.01976051

print("The cross-validation error is minimized for cost = 1")

## [1] "The cross-validation error is minimized for cost = 1"
```

(c)

Now repeat for (b), this time using SVMs with radial and polynomial basis kernels, with different values of gamma and degree and cost. Comment on your results.

```
set.seed(personal)
y = tune(svm, binary.mpg ~., data = Auto, kernel = "polynomial", ranges = list(cost = c(0.01, 0.1, 1, 5, 10, 100), degree = c(2, 3, 4)))
z = tune(svm, binary.mpg ~., data = Auto, kernel = "radial", ranges = list(cost = c(0.01, 0.1, 1, 5, 10, 100), gamma = c(0.01, 0.1, 1, 5, 10, 100, 1000)))
summary(y)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost degree
##   100      2
##
## - best performance: 0.3060897
##
## - Detailed performance results:
##   cost degree   error dispersion
## 1 1e-02      2 0.5535897 0.04171890
## 2 1e-01      2 0.5535897 0.04171890
## 3 1e+00      2 0.5535897 0.04171890
## 4 5e+00      2 0.5535897 0.04171890
## 5 1e+01      2 0.4844231 0.11172253
## 6 1e+02      2 0.3060897 0.05506460
## 7 1e-02      3 0.5535897 0.04171890
## 8 1e-01      3 0.5535897 0.04171890
## 9 1e+00      3 0.5535897 0.04171890
## 10 5e+00     3 0.5535897 0.04171890
## 11 1e+01     3 0.5535897 0.04171890
## 12 1e+02     3 0.3445513 0.06156313
## 13 1e-02     4 0.5535897 0.04171890
## 14 1e-01     4 0.5535897 0.04171890
## 15 1e+00     4 0.5535897 0.04171890
## 16 5e+00     4 0.5535897 0.04171890
## 17 1e+01     4 0.5535897 0.04171890
## 18 1e+02     4 0.5535897 0.04171890

summary(z)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
```

```
##
## - best parameters:
## cost gamma
## 100 0.01
##
## - best performance: 0.01532051
##
## - Detailed performance results:
## cost gamma error dispersion
## 1 1e-02 1e-02 0.57647436 0.03687622
## 2 1e-01 1e-02 0.08923077 0.06276147
## 3 1e+00 1e-02 0.07403846 0.04271928
## 4 5e+00 1e-02 0.05115385 0.04018094
## 5 1e+01 1e-02 0.02557692 0.01709522
## 6 1e+02 1e-02 0.01532051 0.01788871
## 7 1e-02 1e-01 0.19852564 0.08608860
## 8 1e-01 1e-01 0.08166667 0.05510683
## 9 1e+00 1e-01 0.05628205 0.03983401
## 10 5e+00 1e-01 0.02814103 0.01893035
## 11 1e+01 1e-01 0.02044872 0.02020886
## 12 1e+02 1e-01 0.02301282 0.02244393
## 13 1e-02 1e+00 0.57647436 0.03687622
## 14 1e-01 1e+00 0.57647436 0.03687622
## 15 1e+00 1e+00 0.06378205 0.03674375
## 16 5e+00 1e+00 0.06641026 0.03678591
## 17 1e+01 1e+00 0.06641026 0.03678591
## 18 1e+02 1e+00 0.06641026 0.03678591
## 19 1e-02 5e+00 0.57647436 0.03687622
## 20 1e-01 5e+00 0.57647436 0.03687622
## 21 1e+00 5e+00 0.51762821 0.05340278
## 22 5e+00 5e+00 0.51256410 0.06327615
## 23 1e+01 5e+00 0.51256410 0.06327615
## 24 1e+02 5e+00 0.51256410 0.06327615
## 25 1e-02 1e+01 0.57647436 0.03687622
## 26 1e-01 1e+01 0.57647436 0.03687622
## 27 1e+00 1e+01 0.54839744 0.05805619
## 28 5e+00 1e+01 0.53814103 0.05381159
## 29 1e+01 1e+01 0.53814103 0.05381159
## 30 1e+02 1e+01 0.53814103 0.05381159
## 31 1e-02 1e+02 0.57647436 0.03687622
## 32 1e-01 1e+02 0.57647436 0.03687622
## 33 1e+00 1e+02 0.57647436 0.03687622
## 34 5e+00 1e+02 0.57647436 0.03687622
## 35 1e+01 1e+02 0.57647436 0.03687622
## 36 1e+02 1e+02 0.57647436 0.03687622
## 37 1e-02 1e+03 0.57647436 0.03687622
## 38 1e-01 1e+03 0.57647436 0.03687622
## 39 1e+00 1e+03 0.57647436 0.03687622
## 40 5e+00 1e+03 0.57647436 0.03687622
```

```
## 41 1e+01 1e+03 0.57647436 0.03687622
## 42 1e+02 1e+03 0.57647436 0.03687622

print("The lowest cross-validation error for a polynomial kernel, is obtained
for a degree of 2 and a cost of 100.")

## [1] "The lowest cross-validation error for a polynomial kernel, is obtained
d for a degree of 2 and a cost of 100."

print("The lowest cross-validation error for a radial kernel, is obtained for
a gamma for 0.01 and a cost of 100.")

## [1] "The lowest cross-validation error for a radial kernel, is obtained fo
r a gamma for 0.01 and a cost of 100."
```


(d)

Make some plots to back up your assertions in (b) and (c). Hint: In the lab, we used the `plot()` function for svm objects only in cases with $p=2$. When $p>2$, you can use the `plot()` function to create plots displaying pairs of variables at a time. Essentially, instead of typing `plot(svmfit, dat)` where `svmfit` contains your fitted model and `dat` is a data frame containing your data, you can type `plot(svmfit, dat, x1~x4)` in order to plot just the first and fourth variables. However, you must replace `x1` and `x4` with the correct variable names. To find out more, type `?plot.svm`.


```
set.seed(personal)
svm_linear = svm(binary.mpg~., data = Auto, kernel = "linear", cost = 1)
svm_polynomial = svm(binary.mpg~., data = Auto, kernel = "polynomial", cost =
100, degree = 2)
svm_radial = svm(binary.mpg~., data = Auto, kernel = "radial", cost = 100, ga
mma = 0.01)

plottings = function(fitting) {
  for (i in names(Auto)[!(names(Auto) %in% c("mpg", "binary.mpg", "name"))]) {
    plot(fitting, Auto, as.formula(paste("mpg~", i, sep = "")))
  }
}
plottings(svm_linear)
```



SVM classification plot




SVM classification plot




SVM classification plot




SVM classification plot




SVM classification plot



SVM classification plot




SVM classification plot




```
plottings(svm_polynomial)
```


SVM classification plot




SVM classification plot




SVM classification plot




SVM classification plot




SVM classification plot



SVM classification plot




SVM classification plot




```
plottings(svm_radial)
```


SVM classification plot




SVM classification plot




SVM classification plot




SVM classification plot



SVM classification plot



SVM classification plot



SVM classification plot

